

Similar Days at Airports in the New York Area for Air Traffic Flow Management Planning

Akhil Shah, Kenneth Kuhn, Chris Skeels

June 2, 2015

Executive Summary

This report describes our identification of sets of similar days where similarity is defined in terms of the conditions relevant to the planning of various Air Traffic Flow Management Initiatives (ATFMIs) such as Ground Delay Programs (GDPs). The work here represents a first step toward the construction of a decision support tool for dispatchers at airline operations centers and officials at the Federal Aviation Administration’s Air Traffic Control System Command Center (ATCSCC) that plan ATFMIs. A side product of our work is an overview of reasonable approaches for categorizing calendar days in aviation systems research. We describe the identification of similar days using the following sequence of steps: collecting available and appropriate data, defining features within the data, applying clustering or classifying algorithms, and assessing the obtained results.

Terminal Area / Aerodrome Forecast (TAF), Aviation Routine Weather Report (METAR), and Localized Aviation MOS (Model Output Statistics) Product (LAMP) data are all useful data sets for describing weather at airports in the New York area (and elsewhere). TAF and LAMP data describe forecast conditions while METAR data describe observed conditions. Historical archives of all three types of data exist, although coverage and availability are limited when it comes to TAF data. Aviation System Performance Metrics (ASPM) data summarize scheduled and observed airport operations. ATFMI advisory data can be analyzed to see when and where ATFMIs have been implemented.

Reasonable methods for defining features include: applying expert judgment (a “knowledge-based” approach), using Principal Component Analysis to find weighted summations of individual observations that capture variance among days across the observations (a “PCA-based” approach), and summarizing weather conditions using weighted summations of individual observations where scheduled traffic levels determine weights (a “traffic-biasing” approach). We favor the knowledge-based approach here given the many relevant published research reports identifying features.

Clustering is typically considered unsupervised learning, with the goal of exploring and finding structure in feature data, while classifying is considered supervised learning, with the goal of modeling and then forecasting label data. Many efforts in aviation systems use records on the presence or absence of GDPs at various time points as label data. We do not wish to model current decision making. DBSCAN and k-means algorithms are commonly used for clustering. DBSCAN will not assign an outlying data point to a cluster. We prefer k-means clustering instead as we wish to assign all days to a cluster containing the most similar days.

Table 1 details examples of reasonable methodologies for categorizing calendar days for aviation systems research. The final row, highlighted in bold, represents our preferred approach.

Table 1: Reasonable Approaches for Clustering Days by Conditions at Airports

Goal	Raw Data	Feature Selection	Labels	Model / Algorithm
Descriptive model of GDP initiation	TAF data, ASPM data, ATFMI advisory data	Knowledge-based	Initiation / non-initiation of GDP	Logistic regression
Forecast (near future) ATFMI	TAF data, ASPM data, ATFMI advisory data	Knowledge-based	Presence / absence of ATFMI	Random forest
Cluster days according to observed conditions	METAR data, ASPM data, ATFMI advisory data	PCA-based	- (unsupervised)	k-means
Explore structure in airport weather data	METAR data, ASPM data	Traffic-biasing	-	DBSCAN
Cluster days according to forecast conditions	LAMP data, ASPM data	Knowledge-based	-	k-means

We apply our preferred approach, and a few alternative approaches. The results show...

1 Introduction and Context

Personnel at the Federal Aviation Administration’s Air Traffic Control System Command Center and at airline operations centers regularly implement Air Traffic Flow Management Initiatives (ATFMIs) such as Ground Delay Programs (GDPs) purposefully delaying, canceling, and rerouting flights. These initiatives increase the safety and efficiency of the nation’s air transportation system, for example by replacing airborne delay with ground delay, and are necessary during inclement weather and in other situations where demand for system resources exceeds capacity. In particular, problems at airports often create the need for ATFMIs. Analysis of the past use of ATFMIs can demonstrate the relative success of courses of action but must account for the distinct conditions faced during planning and operations. An identification of days that are similar can help, for example allowing analysts to focus on the 10 days in the past two years when there was thunderstorm activity at the key airports in the New York area between 8am and 11am, local time, but clear weather the rest of the day. As one study reported, “clustering techniques appear to be promising methods for identifying the major causes of Ground Delay Programs” ([5]).

This report describes our work to develop methodologies for the identification of similar days in terms of aviation weather and air traffic operations at the airports in the New York area. This report follows an earlier report to identify similar days based on conditions in the airspace around New York City. We do not wish to replicate the prior report and thus only report on new findings specific to our study of airports. The earlier report includes more detail regarding why it would be beneficial to identify similar days from the perspective of ATFMI planning or operations.

As in our earlier work that focused on the airspace, we have published many of our results focusing on airports in a web based application. The application is a minor update of the version we developed and reported on previously. The earlier report contains a more complete description of our web based application, and we here only highlight recent changes.

In this report, we primarily report on our work to collect data, to identify features that describe aviation weather and air traffic at airports in the New York area, and interesting results we obtain when applying well-known clustering algorithms. The final section of this report is a conclusion that compares the airport-focused results to the airspace-focused results obtained previously.

2 Data Collection

We are interested in describing forecast and observed weather and air traffic at airports in the New York area. We focus on John F. Kennedy International Airport (JFK), Newark Liberty International Airport (EWR), and LaGuardia Airport (LGA). These are the busiest airport in the region. Our methods could easily be applied to other airports in the area, or in other areas.

2.1 Airport Weather Data

Airports themselves issue Terminal Area/Aerodrome Forecast reports (TAFs) and Aviation Routine Weather Reports (METARs) which summarize local weather conditions. METARs can contain select forecast data but, generally speaking, TAF data are forecast data while METAR data are observational data. TAF and METAR data contain information on: wind speed and direction, wind gusts, visibility, precipitation, cloud height, cloud cover, humidity, and pressure. TAFs and METARs are issued roughly hourly to ensure reports keep up with changing weather conditions but also that distinct consumers of the data have a consistent report and time to plan against it.

Prior research efforts have linked many of the variables reported in TAF and METAR data to traffic flow management initiatives. [9] used TAF data to forecast Ground Delay Program initiation,

without giving details on the relative importance of specific variables. The authors in [8] point out the relevance of hourly observations of visibility, cloud height, wind, convection, and precipitation in particular, again for predicting GDPs.

A collection of hourly observations of TAF and METAR data, or other data detailing the variables included in TAF and METAR data, covering the busiest airports in the New York area over an extended period of time would comprise an ideal data set to describe airport weather in the area at the time. Current TAF and current METAR data are available at various websites. A large volume of METAR data has been collected in a historical archive that is publicly accessible at <http://www.wunderground.com/history>. We’ve written a script to collect large volumes of this METAR data. While there are repositories of TAF data, such as www.ogimet.com and the NASA Data Warehouse described in [3], none have the same coverage and availability as the repository of METAR data hosted by [wunderground.com](http://www.wunderground.com). This result is unsurprising; there are relatively many uses for historical observations of weather conditions and relatively few for historical forecasts. There is another source of forecast airport weather data: the National Oceanic and Atmospheric Administration (NOAA) Localized Aviation MOS (Model Output Statistics) Product (LAMP). A large historical archive of this data is free to download from NOAA.

An overview of LAMP can be found in [4]. LAMP data is produced roughly hourly and covers airports in the continental United States, Hawaii, Alaska and Puerto Rico. LAMP uses both METAR and radar data, 16-level 2-km radar data from Weather Science, Inc. and 7-level 10-km Radar Coded Message Mosaic data, to forecast statistics that important in aviation systems. Radar data are, in particular, processed to yield predictors of thunderstorm development. Observed lightning from the National Lightning Detection Network are also used in the development of thunderstorm probabilities.

The following are forecast by LAMP: the probability of precipitation, the precipitation type (liquid rain, snow, or freezing rain) conditional on precipitation occurring, the probabilities of cloud ceiling heights belonging to different ranges, the probabilities of total sky cover belonging to different categories (e.g. clear, few, overcast, etc.), the probabilities of visibility belonging to different ranges, the probabilities of obstruction to vision belonging to different categories (e.g. haze, mist, fog, etc.), and the probabilities of thunderstorms in different 2 hour windows of time up to 25 hours into the future.

2.2 Airport Traffic and Other Data

Traffic flow management is concerned with mitigating temporary supply and demand imbalances. At the level of an airport, the primary concern is almost always that the number of aircraft scheduled to land at and take off from a set of runways across a set block of time may be higher than the throughput of the runways will allow. A reasonable length for such a block of time would be an hour or a few hours. Supply-demand imbalances over shorter periods of time, e.g., two aircraft scheduled to land at the same time, can be accommodated with minor path adjustments rather than the more strategic traffic flow management initiatives. Hourly observations of scheduled operations could be easily compared to weather data such as TAFs and METARs that are reported hourly. A collection of hourly observations of the numbers of aircraft scheduled to land at and take off from the busiest airports in the New York area would be an ideal data set here. Aviation System Performance Metrics (ASPM) data includes exactly such data and is available to all.

Table 2 summarizes the availability of useful airport data.

Table 2: Summary of Relevant Airport Data

Acronym	Name	Notes
TAF	Terminal Area/ Aerodrome Forecast	Hourly summaries of forecast weather. Gaps exist in available historical record.
METAR	Aviation Routine Weather Report	Hourly summaries of observed weather. Comprehensive historical archive at wunderground.com.
LAMP	Localized Aviation MOS Product	Hourly forecasts of future weather conditions. Comprehensive historical archive available from NOAA.
ASPM	FAA Aviation System Performance Metrics	Hourly counts of scheduled and observed arrivals and departures. Comprehensive historical archive available from FAA.
ASDI	Aircraft Situation Display to Industry	Detailed aircraft flight plan and observed position data. Comprehensive historical archive available from NASA.

3 Selecting Features to Describe Airport Conditions

It is important to note that the dominant factor which determines the quality of any machine learning (or more generally, statistical model) approach is the selection of features that are most relevant rather than the choice of prediction or clustering algorithms. Generally, the best source for relevant features is expert domain knowledge. However machine learning methods for automated feature selection can complement expert elicitation by quantifying the relevance of features to study goals such as modeling specific outcomes variables.

Feature selection methods often rely on this idea that the goal of a study to model specific categorical or continuous variables and that a dataset with accompanying ground truth **label data** is provided. This is the **supervised learning** case. For example, if we were interested in determining which weather and traffic features were most relevant to predicting a Ground Delay Program (GDP) on any given day, the feature selection algorithm would require a **training dataset** that included not only the predictors (weather and traffic) but also the observed ground-truth label of absence or presence of GDP for each data record [8]. The goals of this algorithm would be to model GDP decision making, using the training dataset to set the model, so that the derived model is able to forecast GDP decision making in a separate **test dataset** as well as in subsequently considered datasets.

Our primary goal is unusual; we wish to explore, characterize, and categorize airport weather and traffic conditions without modeling current GDP decision making. Therefore, it was not appropriate, for our study, to use ground-truth labels. The lack of ground-truth labels rules out various feature selection methods [6]. Methods for feature selection and cluster analysis in an **unsupervised learning** case like ours are often based on finding the most clearly defined clusters possible. While we are interested in finding clearly defined clusters, where possible, we are most interested in identifying features that summarize and characterize how analysts and air traffic control planners describe airport weather and air traffic at airports. We thus report on measures of **structure** in our feature data, particularly as a function of the parameters of clustering algorithms that cannot be set using expert judgement, but do not believe such measures should be exclusively used to define the features we use or to ascertain the relative importance of these features.

Our clustering results were derived from features that were selected by either surveying the

literature for those considered relevant or using heuristics. In the sections below we summarize the literature employed for knowledge-based feature selection and also our heuristic traffic-biasing and PCA-based methodologies.

3.1 Knowledge-Based Feature Selection

Previous studies have employed airport-centric observations and forecasts of weather and traffic data for clustering days. We will use these weather and traffic features identified as relevant as part of our overall clustering analysis. In these studies, features are initially identified based on domain knowledge, and then filtered for relevancy using either correlation analysis or other statistical measures, when used as explanatory variables in a statistical model for predicting observed ATFMI data [8, 5].

In [8], the authors employ logistic regression and decision trees to predict the absence or presence of GDP every hour at EWR. Explanatory variables which are determined as statistically significant as predictors of GDPs at EWR include: visibility, runway crosswind, demand-to-capacity ratio and queueing delay as derived from hourly scheduled arrival data and the assumed arrival capacity, Weather Impacted Traffic Index for the local Air Route Traffic Control Center (ZNY - New York Center), and averages, over the preceding three hours, of wind speed, demand-to-capacity ratio, and ZNY WITI.

In [5], the authors performed clustering of EWR airport-level data which included observed hourly arrivals and hourly wind speed, wind gusts, wind direction, and ceiling. These features were selected as relevant by performing a correlation analysis with the absence or presence of hourly GDP at EWR. Note unlike the previous analysis, visibility at EWR was determined to be not as strong a predictor of GDP, and thus was not used in subsequent clustering.

In [7], the authors also attempted to model GDP initiation at EWR. Weather variables considered included: the log of cloud ceiling heights, the log of (categorical) visibility data, and binary variables describing whether it was snowing or not, whether there was thunderstorm activity or not, whether tailwinds were greater than 5 knots or not, whether crosswinds were greater than 10 knots or not, whether Instrument Meteorological Conditions (IMC) or Visual Meteorological Conditions (VMC) were in effect. Scheduled arrival demand and airport capacity estimates were also used as explanatory variables. Logistic regression and similarity-based logistic regression models were used to forecast GDPs using data from 1, 2, or 3 hours prior. At the same conference, the authors of [1] described applied behavioral cloning and cascaded supervised inverse reinforcement learning to model GDP initiation. In addition to the data described above, these authors used runway configuration and delay data.

In both of the studies cited in the preceding paragraph, the authors pointed out that the relative infrequency of GDP initiation in their training datasets hampered modeling efforts. A related study [10] analyzed combinations of features (“queries”) from airport-specific Weather Impacted Traffic Index (WITI) data to determine which were most relevant in predicting observed weather-related GDP from present time to six hours in the future. They developed an information-retrieval model consisting of thresholding WITI features as queries or rules. Their model was trained using observed 2008 GDP data and tested on 2009 GDP data. However this study concluded that the rules based on WITI features were only slightly more predictive than a constant rule which always predicted “no GDP.” According to results discussed by [7], the overall accuracy of one model was 90% but, when focusing on the observed, actual initiation of GDPs, the model only successfully forecast initiation roughly 30% of the time.

In [2], the authors modeled airport capacity profiles at EWR, San Francisco International Airport (SFO), and Los Angeles International Airport (LAX). The authors used TAF numerical data on

wind speed and direction, visibility, and cloud heights, and TAF categorical data on the presence or absence of rain, fog, and mist. They tried k-means clustering on all of the 420 data points that describe a day in their dataset in one experiment, applied a PCA-based approach similar to the one we describe later in this chapter to automatically generate features in another experiment, and use a “dynamic time warping” approach that attempts to model time shift effects where changes in weather at one time point can lead to changes in airport capacity at another time point in a third experiment.

As an alternative to explicitly modeling the presence or absence of GDPs, the authors of [11] employ a semi-supervised clustering methodology. One of the main contributions in the cited study is the derivation of a distance metric for quantifying the separation between distance data points that describe hours or days from user specified sets of similar and dissimilar data points. The authors also suggest and apply a specific method for defining sets of similar and dissimilar hours based on runway configuration, Airport Acceptance Rates (AAR), and whether visual or instrument meteorological conditions are in effect. The derived (learned from data and user supplied specifications) distance metric minimizes the distance between weather feature vectors at “similar” hours, while ensuring the distance between feature vectors at “dissimilar” hours is greater than one. The authors’ proof-of-concept example uses historical TAF and ASPM data from 2011 to define a distance metric and, for two given days in 2012 at EWR, selects the five most similar days from 2011. Their feature vector is nine-dimensional and consists of boolean values for absence/presence of thunderstorm and snow, and various functional transformations of visibility, ceiling and wind speed.

Based on the literature, we define features that include, for each day, three-hour averages for forecast data on each of the following variables: scheduled arrivals, scheduled departures, visibility, and wind speed.

3.2 A Traffic-Biasing Approach

Both our weather and traffic datasets contain hourly observations of many variables. Attempting to use all observations as features would produce a data set that is too large for most clustering algorithms (due to the curse of dimensionality). A naïve mitigation strategy would be to uniformly average the values of numerical variables over a set of 24 hourly observations to produce a single value for a given day. This would, however, treat all observations as equally important. We are interested in weather as it relates to air traffic and we know that not all times of day are equally important when it comes to air traffic. A more sophisticated approach would be to employ a weighted average approach where weights are determined through the following heuristic: determine daily average arrival count for a given airport, and normalize these counts to sum to one. In our case we used ASPM scheduled arrival and departure counts from 2010-2013 to determine the average hourly arrival for each of the three New York region airports in our analysis. Once normalized, these form 24 weights w_h , where h is the hour index, which are then used to calculate weighted averages of METAR numerical variables. For example, a given day’s visibility at JFK would have the following expression to reduce the 24 observed values to a single value $v_d = \sum_{h=1}^{24} w_h v_{d,h}$, where d is the index for a given day in the dataset. We call this approach traffic biasing.

Figure 1 shows the hourly variation of arrival data at JFK for every day between 2010-2013. Note that there is a clear pattern; some hours are much busier than others. The traffic biasing heuristic uses this information, weighting weather observations more when they were recorded at a time when there was more air traffic scheduled to use the airport.




Figure 1: Counts of scheduled arrivals at JFK by hour.

3.3 A PCA-Based Approach

Another heuristic we have explored for feature selection is based on Principal Component Analysis (PCA) and, similar to traffic biasing, converts hourly observations of a variable to a daily summary statistic. PCA is used to select linear combinations of the 24 observed values of each variable per day which capture the largest possible variation between days (i.e. the first n principal components). Our earlier airspace-focused report describes, in more detail, how PCA can be used to automatically select features for a subsequent cluster analysis. We summarize that discussion by noting that our application of PCA is a way to reduce the dimensions of a dataset while maintaining as much of the differences between data points (here days) as possible.

Figure 2 shows the amount of variance explained by the first three principal components of a given day at JFK for various weather variables recorded in METAR data. Note that some variables do have strongly explanatory first principal component, but others do not.




Figure 2: Variance explained by principal components for weather variables at JFK.

4 Results and Discussion

4.1 Exploratory Analysis

Below we summarize the statistical variation of the datasets we have used to derive features and cluster days.

Variation over 24 hours of selected METAR weather variables for days in 2010-2013.

4.2 Assessing Clusters

When ground-truth labels are available, clusters can be assessed by various quantitative metrics that characterize the intra-cluster similarity of those labels versus the inter-cluster similarity of labels. One example of such a metric is the RAND index, but there are various others. More generally a “confusion matrix” is a standard representation in Machine Learning to present the accuracy of an algorithm that models a categorical variable. Typically, a model that has been fit using a training data set is asked to forecast the categorical variable in a distinct test data set. The rows of the confusion matrix represent the actual values of the categorical variable in the test data set while the columns represent the model output values of the categorical variable using the model inputs from the test data set. Each cell in the matrix is a count showing how many times the model predicted label x (column) when the actual observation was label y (row).

We propose a novel method to use categorical data as “psuedo-labels” for each observation. The data we use for these psuedo-labels are the observed “Condition” data column of METAR data, a categorical value. Each day (sample of our data-matrix) has 24 such values and we propose three types of concurrency tables which allow use to analyze the similarity within and between the clusters using the frequency of occurrence of those categorical values.

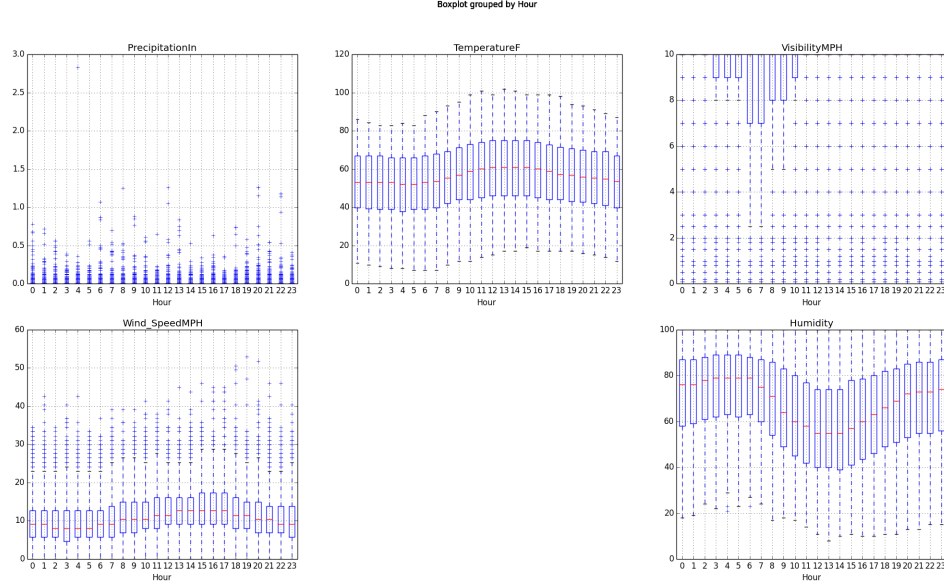


Figure 3: Monthly variation in observations of weather variables

Concurrency Tables

The Condition data within METAR data describe weather conditions at an airport using a categorical variable. The data are not ground-truth labels in the sense that we would explicitly seek to build a model of this data or claim that the data adequately summarize airport weather data for air traffic flow management initiative planning. However, it is worthwhile to compare the results of cluster analysis to the Condition data, to see if, for example, days with the Condition label of “Thunderstorm” appear in one cluster while days with the Condition label of “Clear” appear in another cluster. To explore the relationship between the METAR Condition data and identified clusters, we created concurrency tables. The idea is similar to that of the Confusion Matrix described above. Each row in a concurrency table refers to a specific Condition label and each column to a specific cluster output by a clustering algorithm. Each cell in a concurrency table contains a count of the number of times a day was assigned to a specific Condition label (row) according to observed METAR data and a specific cluster (column) according to the results of a specific form of cluster analysis.

	1	2	3	4	5
Blowing Snow	-	-	-	2	2
Clear	490	548	1936	2651	2674
Fog	5	230	233	238	462
Haze	55	181	202	207	310
Heavy Rain	22	74	80	155	238
Heavy Snow	-	4	4	6	24
Heavy Thunderstorms and Rain	36	62	66	71	86
Ice Pellets	-	3	3	5	6
Light Drizzle	9	144	170	191	458
Light Freezing Drizzle	-	7	7	7	45
Light Freezing Rain	-	4	4	4	36
Light Ice Pellets	-	3	4	4	24
Light Rain	347	874	1330	1867	2799
Light Rain Showers	-	-	-	1	1
Light Snow	-	194	304	463	814
Light Thunderstorms and Rain	70	100	123	128	146
Light Thunderstorms and Snow	-	-	-	-	1
Mist	-	12	12	12	27
Mostly Cloudy	3410	4483	8639	11069	11501
Overcast	1135	2319	4619	5825	7271
Partly Cloudy	950	1075	2924	4404	4455
Patches of Fog	-	4	4	4	5
Rain	40	150	184	264	481
Scattered Clouds	2068	2448	4859	6562	6659
Shallow Fog	2	2	2	2	2
Smoke	-	-	1	1	1
Snow	-	15	18	40	71
Squalls	-	-	1	1	1
Thunderstorm	47	51	62	67	68
Thunderstorms and Rain	24	35	40	47	53
Thunderstorms with Small Hail	-	1	1	1	1
Unknown	5	7	11	16	20

Day-Aggregated counts:

Day-Most frequent counts:

	1	2	3	4	5
Clear	117	117	655	899	899
Fog	-	72	72	72	132
Haze	-	10	10	10	33
Heavy Rain	-	-	-	25	25
Light Drizzle	-	20	20	20	54
Light Freezing Drizzle	-	-	-	-	23
Light Rain	48	282	397	566	1112
Light Snow	-	147	196	262	551
Mostly Cloudy	2616	3205	6215	7825	7903
Overcast	453	1181	2337	2885	3837
Partly Cloudy	334	358	1146	1837	1837
Rain	-	9	9	9	9
Scattered Clouds	1113	1186	2126	2828	2852
Snow	-	-	-	12	22

Day-Least frequent counts:

	1	2	3	4	5
Clear	73	79	252	343	348
Fog	-	12	12	13	17
Haze	11	22	35	38	48
Heavy Rain	4	22	25	35	60
Heavy Snow	-	4	4	4	4
Heavy Thunderstorms and Rain	5	7	7	9	14
Ice Pellets	-	-	-	-	1
Light Drizzle	1	14	23	30	66
Light Freezing Drizzle	-	-	-	-	2
Light Freezing Rain	-	-	-	-	1
Light Ice Pellets	-	1	2	2	7
Light Rain	50	76	150	205	220
Light Rain Showers	-	-	-	1	1
Light Snow	-	3	6	26	27
Light Thunderstorms and Rain	18	22	28	30	30
Light Thunderstorms and Snow	-	-	-	-	1
Mist	-	10	10	10	14
Mostly Cloudy	141	158	366	508	534
Overcast	94	109	281	383	397
Partly Cloudy	189	206	486	661	665
Patches of Fog	-	-	-	-	1
Rain	15	29	40	60	90
Scattered Clouds	178	198	603	752	769
Shallow Fog	2	2	2	2	2
Smoke	-	-	1	1	1
Snow	-	-	-	-	3
Squalls	-	-	1	1	1
Thunderstorm	3	5	5	5	5
Thunderstorms and Rain	10	18	19	22	27
Unknown	1	2	4	4	6

Condition Concurrency Tables for airport:EWR

	1	2	3	4	5
Blowing Snow	-	-	10	10	10
Clear	516	1446	1514	1546	3176
Fog	-	-	70	147	149
Haze	29	35	191	416	525
Heavy Ice Pellets	-	-	-	1	1
Heavy Rain	5	20	90	242	284
Heavy Snow	-	1	18	53	53
Heavy Thunderstorms and Rain	3	15	31	56	128
Ice Pellets	-	-	3	12	12
Light Drizzle	4	25	158	618	656
Light Freezing Drizzle	-	-	7	48	48
Light Freezing Rain	-	-	17	52	52
Light Ice Pellets	3	8	14	39	41
Day-Aggregated counts: Light Rain	151	547	1217	2633	3438
Light Rain Showers	-	-	-	1	1
Light Snow	45	143	434	834	881
Light Thunderstorms and Ice Pellets	-	-	-	2	2
Light Thunderstorms and Rain	3	32	54	88	199
Mist	-	-	-	3	3
Mostly Cloudy	1634	4074	5015	5404	11420
Overcast	886	2021	3168	4695	7838
Partly Cloudy	551	1833	1947	1988	3679
Rain	4	60	182	492	560
Scattered Clouds	881	2537	2835	2932	5746
Smoke	-	-	-	-	1
Snow	6	11	41	91	91
Thunderstorm	5	27	44	47	121
Thunderstorms and Rain	4	17	30	55	94
Unknown	4	12	14	15	17

		1	2	3	4	5		
Day-Most frequent counts:	Blowing Snow	-	-	9	9	9	Day-Least	
	Clear	225	613	613	613	1235		
	Fog	-	-	10	32	32		
	Haze	8	8	40	118	137		
	Heavy Rain	-	-	15	15	15		
	Heavy Snow	-	-	-	17	17		
	Light Drizzle	-	-	11	152	152		
	Light Freezing Drizzle	-	-	-	25	25		
	Light Freezing Rain	-	-	10	10	10		
	Light Rain	17	146	485	1457	1594		
	Light Snow	14	51	295	627	639		
	Mostly Cloudy	1113	2869	3367	3437	7909		
	Overcast	510	953	1627	2476	4220		
	Partly Cloudy	199	670	670	670	1222		
	Rain	-	-	-	44	44		
	Scattered Clouds	362	1017	1083	1083	2154		
		1	2	3	4	5		
frequent counts:	Blowing Snow	-	-	1	1	1	Condition Con-	
	Clear	55	144	145	153	371		
	Fog	-	-	8	12	12		
	Haze	5	8	31	41	58		
	Heavy Ice Pellets	-	-	-	1	1		
	Heavy Rain	1	8	15	56	63		
	Heavy Snow	-	1	2	3	3		
	Heavy Thunderstorms and Rain	-	-	1	4	7		
	Light Drizzle	1	8	24	53	66		
	Light Freezing Drizzle	-	-	1	2	2		
	Light Ice Pellets	-	-	1	4	4		
	Light Rain	14	56	75	105	191		
	Light Rain Showers	-	-	-	1	1		
	Light Snow	12	28	46	55	61		
	Light Thunderstorms and Rain	-	7	12	16	25		
	Mist	-	-	-	1	1		
	Mostly Cloudy	95	233	247	266	511		
	Overcast	98	191	204	253	519		
	Partly Cloudy	96	272	290	299	614		
	Rain	2	13	28	68	90		
	Scattered Clouds	126	246	283	302	686		
	Smoke	-	-	-	-	1		
	Snow	-	-	2	9	9		
	Thunderstorm	-	3	4	5	17		
	Thunderstorms and Rain	2	5	10	13	34		
	Unknown	-	7	8	8	10		

currency Tables for airport:LGA

		1	2	3	4	5
Day-Aggregated counts:	Blowing Snow	-	-	-	1	1
	Clear	605	2618	2630	3566	3618
	Drizzle	-	-	1	1	1
	Fog	5	5	17	19	173
	Haze	40	133	183	193	327
	Heavy Rain	19	65	226	276	354
	Heavy Snow	-	-	21	21	26
	Heavy Thunderstorms and Rain	2	52	84	86	91
	Ice Pellets	-	-	-	-	3
	Light Drizzle	19	58	225	255	374
	Light Freezing Drizzle	3	3	8	8	25
	Light Freezing Rain	1	1	25	25	66
	Light Ice Pellets	3	9	19	20	21
	Light Rain	176	956	2143	2565	3161
	Light Rain Showers	-	1	1	1	1
	Light Snow	115	172	577	703	955
	Light Thunderstorms and Rain	8	87	110	116	129
	Light Thunderstorms and Snow	-	-	2	2	2
	Mist	3	3	3	3	26
	Mostly Cloudy	1823	7653	7940	9819	10413
	Overcast	1084	4695	5847	7214	8654
	Partly Cloudy	632	2851	2883	4185	4262
	Rain	14	91	418	483	622
	Scattered Clouds	964	4073	4131	5424	5623
	Snow	6	6	54	71	106
	Thunderstorm	6	72	73	75	78
	Thunderstorms and Ice Pellets	-	-	1	1	1
	Thunderstorms and Rain	8	32	49	56	57
	Unknown	-	2	2	4	4
		1	2	3	4	5
Day-Most frequent counts:	Clear	271	1253	1253	1653	1666
	Haze	-	-	-	-	34
	Light Drizzle	-	-	41	41	48
	Light Freezing Rain	-	-	13	13	44
	Light Rain	27	190	947	1031	1297
	Light Snow	54	65	412	470	671
	Mostly Cloudy	1299	5344	5414	6465	6651
	Overcast	605	2839	3512	4318	5367
	Partly Cloudy	238	1010	1010	1655	1655
	Rain	-	9	48	48	63
	Scattered Clouds	389	1395	1395	1885	1925

	1	2	3	4	5
Blowing Snow	-	-	-	1	1
Clear	68	354	356	459	466
Drizzle	-	-	1	1	1
Fog	1	1	5	5	6
Haze	6	22	31	31	55
Heavy Rain	1	11	29	38	45
Heavy Snow	-	-	8	8	8
Heavy Thunderstorms and Rain	-	4	6	7	8
Light Drizzle	4	12	34	44	72
Light Freezing Drizzle	-	-	1	1	4
Light Freezing Rain	1	1	1	1	1
Light Ice Pellets	1	2	3	4	5
Light Rain	19	152	187	244	254
Light Rain Showers	-	1	1	1	1
Light Snow	10	21	22	28	29
Light Thunderstorms and Rain	3	21	22	22	27
Mist	3	3	3	3	11
Mostly Cloudy	107	309	329	441	454
Overcast	66	266	280	362	367
Partly Cloudy	115	453	461	571	579
Rain	4	25	46	58	67
Scattered Clouds	133	557	567	687	718
Snow	-	-	2	4	6
Thunderstorm	3	12	12	12	12
Thunderstorms and Rain	4	12	21	21	21

Conclusion

In conclusion,

References

- [1] M. Bloem and N. Bambos. Ground delay program analytics with behavioral cloning and inverse reinforcement learning. In *INFORMS*, 2014.
- [2] G. Buxi and M. Hansen. Generating probabilistic capacity profiles from weather forecast: a design-of-experiment approach. *Berlin, Germany*, 2011.
- [3] M. Eshow, M. Lui, and S. Ranjan. Architecture and capabilities of a data warehouse for atm research. In *Digital Avionics Systems Conference*, 2014.
- [4] J. Ghirardelli. An overview of the redeveloped localized aviation mos program (lamp) for short-range forecasting. In *American Meteorological Society*, 2005.
- [5] S. Grabbe, B. Sridhar, and A. Mukherjee. Similar days in the nas: an airport perspective. In *AIAA Aviation Technology, Integration, and Operations (ATIO) Conference*, 2013.

- [6] I. Guyon and A. Elisseeff. An introduction to variable and feature selection. *The Journal of Machine Learning Research*, 3:1157–1182, 2003.
- [7] M. Hansen, Y. Liu, M. Seelhorst, M. Ying, and A. Pozdnukhov. Predicting the initiation of a ground delay program. In *INFORMS*, 2014.
- [8] A. Mukherjee, S. Grabbe, and B. Sridhar. Predicting ground delay program at an airport based on meteorological conditions.
- [9] D. Smith and L. Sherry. Decision support tool for predicting ground delay programs (gdp) and airport delays from weather forecast data. In *Annual Transportation Research Board (TRB) Conference, Washington, DC*, 2009.
- [10] S. R. Wolfe and J. L. Rios. A method for using historical ground delay programs to inform day-of-operations programs. In *Proc. of AIAA Guidance, Navigation, and Control Conference*, 2011.
- [11] L. Yi, M. Hansen, A. Pozdnukhov, and M. Ball. Assessing terminal weather forecast similarity for strategic air traffic management. *Sixth International Conference on Research in Air Transportation*, May 2014.